


Computing Power: HPC Tutorial Session



Bocconi Students for Machine Learning (BSML)

April 16, 2025

Introduction

In this tutorial, we go through an example use of the high performance computing (HPC) cluster. We will cover:

1. Login using ssh
2. VPN usage
3. Quality of life improvements (ssh keys, ssh configuration)
4. Create & submit a job
5. Misc commands

1 Login Using SSH

Connect to bocconi-students Wi-Fi. Open terminal and ssh into students HPC using password:

```
[local]$ ssh bocconi_id@slnode-da.sm.unibocconi.it
bocconi_id@slnode-da.sm.unibocconi.it's password:
```

Close connection:

```
[hpc]$ logout
```

2 VPN Usage

Trying to access the cluster from a different Wi-Fi network causes an error:

```
[local]$ ssh bocconi_id@slnode-da.sm.unibocconi.it
ssh: Could not resolve hostname slnode-da.sm.unibocconi.it: No such host is known.
```

You need to first setup and connect to the FortiClient VPN.

3 Quality of Life (SSH Keys, SSH Config)

3.1 SSH Keys

We want to setup ssh keys to avoid typing the password every time.

We need to have a public ssh key, call it `id_rsa.pub` copied inside the file `.ssh/authorized_keys` on the HPC. The corresponding private key, call it `id_rsa`, must be stored inside the `.ssh/` directory on the local machine.

Whenever we try to access the HPC, the ssh protocol will send a test message encrypted with the public key. Only if we have the correct private key we are able to decode it and send it back, allowing for secure access.

We create a pair of keys using the `ssh-keygen` command from inside the local directory `.ssh/` (add a signature to recognize the key):

```
[local]$ cd .ssh
[local]$ ssh-keygen -C "my-new-key"
```

You can either change the file name or keep the default `id_rsa`.

Now print the content of the public key and copy it:



```
[local]$ type id_rsa.pub
ssh-ed25519 AAAAC3NzaC1lZDI1NTE5AAAAII0kb/JST6a//zakM6CMT1UKbVk7FwXWu75pzH/YfP9B
my-new-key
```

Login to the HPC and paste it in `.ssh/authorized_keys` (create the file if it does not exists):

```
[local]$ ssh bocconi_id@slnode-da.sm.unibocconi.it
bocconi_id@slnode-da.sm.unibocconi.it's password:
[hpc]$ cd .ssh
[hpc]$ nano authorized_keys
```

Paste the public key and save. Logout and try connecting again (it should work without password).

```
[hpc]$ logout
[local]$ ssh bocconi_id@slnode-da.sm.unibocconi.it
```

3.2 SSH Config

To avoid typing the entire user and host name at every ssh login, open `C:/Users/name/.ssh/config` and add aliases:

```
Host hpc_stud
  HostName slnode-da.sm.unibocconi.it
  User bocconi_id
```

Connect to hpc again:

```
[local]$ ssh hpc_stud
[hpc]$
```

At this point, you should be able to login without using the password and without typing the entire user and host name. We are now ready to create and submit our first SLURM job.

4 Create & Submit a SLURM Job

In this section, we will walk you through all the steps needed to submit the first job from inside the HPC. First, we create an anaconda environment. Then, we create a SLURM job file to execute a demo python script. Finally, we submit the job and monitor the runtime.

4.1 Create Conda Environment

Create a new anaconda environment:

```
[hpc]$ conda create -n bsml python==3.12
```

If needed, install the required dependencies:

```
[hpc]$ conda activate bsml
[hpc]$ python -m pip install torch
...
```

4.2 Create a Python Script

Create a new `.py` file using the Linux command `nano`:

```
[hpc]$ nano bsml.py
```



An interactive text editor will pop-up. Add your code inside:

```
import time
for i in range(10):
    print(f"{i} Connecting Minds, Not Only Layers")
    time.sleep(2)
```

4.3 Create a SLURM Job Script

Again, use `nano` to create a `.sh` file:

```
[hpc]$ nano bsml.sh
```

Fill it using the SLURM syntax:

```
#!/bin/bash
#SBATCH --job-name="bsml"
#SBATCH --account=bocconi_id
#SBATCH --partition=stud

#SBATCH --cpus-per-task=1
#SBATCH --gpus=1
#SBATCH --mem=10GB

#SBATCH --chdir=.
#SBATCH --output=/home/bocconi_id/out.out
#SBATCH --error=/home/bocconi_id/err.err

#SBATCH --mail-type=ALL
#SBATCH --mail-user=giacomo.ciro@unibocconi.it

module load modules/miniconda3
eval "$(conda shell.bash hook)"

conda activate bsml

python bsml.py

module unload modules/miniconda3
echo "The end"
```

4.4 Submit & Monitor the Job

Submit the job:

```
[hpc]$ sbatch bsml.sh
```

Monitor the status of the job:

```
[hpc]$ squeue
```

Cancel the job:

```
[hpc]$ scancel <job_id>
```

5 Misc Commands

We list some SLURM and Linux commands which might come handy when working on the HPC.



5.1 SLURM

List of past jobs

```
[hpc]$ sacct
```

Nodes info:

```
[hpc]$ sinfo
```

More node-specific info:

```
[hpc]$ scontrol show node <node_name>
```

Partition-specific info:

```
[hpc]$ scontrol show partition <partition_name>
```

5.2 Linux

List content of a directory (omit path to list current dir):

```
[hpc]$ ls path/to/dir
```

Print content of file:

```
[hpc]$ cat path/to/file
```

Remove file:

```
[hpc]$ rm path/to/file
```

Remove entire directory recursively:

```
[hpc]$ rm -r path/to/dir
```

Copy file:

```
[hpc]$ cp path/to/file_to_copy path/to/new_file_name
```

Print end of file and refresh when new changes occur:

```
[hpc]$ tail -f path/to/file
```

Run a command every 1 second:

```
[hpc]$ watch -n 1 <cmd>
```

Memory info:

```
[hpc]$ free -h
```

CPU info:

```
[hpc]$ lscpu
```

GPU info:

```
[hpc]$ nvidia-smi
```



5.3 Transfer Files

From local machine to HPC:

```
[local]$ scp /path/to/local/file bocconi_id@node-da.sm.unibocconi.it:/path/to/  
remote/directory
```

From HPC to local machine

```
[local]$ scp bocconi_id@node-da.sm.unibocconi.it:/path/to/remote/file /path/to/  
local/directory
```

Use the flag `-r` to copy a directory recursively.

